

In [150]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

In [151]:

```
df = pd.read_csv('C:/TreeVista/Test 1.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
customer_id      10000 non-null int64
demographic_slice 10000 non-null object
country_reg      10000 non-null object
ad_exp           10000 non-null object
est_income       10000 non-null float64
hold_bal         10000 non-null float64
pref_cust_prob   10000 non-null float64
imp_cscore       10000 non-null int64
RiskScore        10000 non-null float64
imp_crediteval   10000 non-null float64
axio_score       10000 non-null float64
card_offer       10000 non-null bool
dtypes: bool(1), float64(6), int64(2), object(3)
memory usage: 869.2+ KB
```

In [152]:

```
df.columns
```

Out[152]:

```
Index(['customer_id', 'demographic_slice', 'country_reg', 'ad_exp',
       'est_income', 'hold_bal', 'pref_cust_prob', 'imp_cscore', 'RiskScore',
       'imp_crediteval', 'axio_score', 'card_offer'],
      dtype='object')
```

In [153]:

```
df.describe()
```

Out[153]:

| | customer_id | est_income | hold_bal | pref_cust_prob | imp_cscore | RiskScore | imp_crediteval | axio_s |
|-------|---------------|---------------|--------------|----------------|--------------|--------------|----------------|-----------|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.00 |
| mean | 496819.831400 | 65853.355259 | 20.962621 | 0.329419 | 662.548800 | 670.042869 | 25.692162 | 0.393211 |
| std | 287391.314157 | 31093.369592 | 18.841121 | 0.223299 | 90.549985 | 89.965854 | 1.889274 | 0.288243 |
| min | 244.000000 | 2.054543 | -2.140206 | 0.001781 | 500.000000 | 324.436647 | 21.363123 | -0.000050 |
| 25% | 245172.500000 | 39165.786086 | 6.150577 | 0.156965 | 600.000000 | 609.231181 | 24.295435 | 0.139424 |
| 50% | 495734.000000 | 76903.628763 | 11.913366 | 0.272263 | 655.000000 | 669.493442 | 25.611903 | 0.337841 |
| 75% | 745475.250000 | 91032.514900 | 32.238914 | 0.459890 | 727.000000 | 730.484985 | 27.062519 | 0.624886 |
| max | 999870.000000 | 150538.809704 | 81.759632 | 1.144357 | 849.000000 | 1004.497869 | 30.131214 | 1.000000 |

Cleaning the data

In [154]:

```
df.isna().any()
```

Out[154]:

```
customer_id      False
demographic_slice False
country_reg      False
ad_exp           False
est_income       False
hold_bal         False
pref_cust_prob   False
imp_cscore       False
RiskScore        False
imp_crediteval   False
axio_score       False
card_offer       False
dtype: bool
```

In [155]:

```
# Create a correlation matrix
corr_metrics = df.corr()
corr_metrics.style.background_gradient()
```

Out[155]:

| | customer_id | est_income | hold_bal | pref_cust_prob | imp_cscore | RiskScore | imp_crediteval | axio_score |
|----------------|-------------|--------------|-------------|----------------|-------------|-------------|----------------|-------------|
| customer_id | 1 | 0.00492495 | 0.00685613 | -0.00571561 | 0.00910654 | -0.00420723 | 0.0056651 | 0.000761587 |
| est_income | 0.00492495 | 1 | 0.0103311 | 0.00868931 | 0.00351446 | 0.00453002 | -0.000173806 | -0.00467278 |
| hold_bal | 0.00685613 | 0.0103311 | 1 | 0.00182511 | 0.269361 | 0.0139309 | 0.256165 | 0.000158183 |
| pref_cust_prob | -0.00571561 | 0.00868931 | 0.00182511 | 1 | -0.0114992 | -0.0125698 | -0.013733 | -0.0203884 |
| imp_cscore | 0.00910654 | 0.00351446 | 0.269361 | -0.0114992 | 1 | -0.00480859 | 0.926908 | 0.00536186 |
| RiskScore | -0.00420723 | 0.00453002 | 0.0139309 | -0.0125698 | -0.00480859 | 1 | -0.00435891 | 0.000654177 |
| imp_crediteval | 0.0056651 | -0.000173806 | 0.256165 | -0.013733 | 0.926908 | -0.00435891 | 1 | 0.00672548 |
| axio_score | 0.000761587 | -0.00467278 | 0.000158183 | -0.0203884 | 0.00536186 | 0.000654177 | 0.00672548 | 1 |
| card_offer | -0.00680401 | 0.27753 | 0.0561434 | 0.630311 | 0.0228125 | -0.00360271 | 0.0170552 | -0.0153 |

In [158]:

```
# it seems feature 'imp_cscore' is correlated to 'imp_crediteval' , it's good to drop one of them
for the sake of curse of Dim..

df_updated= df.drop(columns=['imp_crediteval'])
```

In [159]:

```
df_updated.describe()
```

Out[159]:

| | customer_id | est_income | hold_bal | pref_cust_prob | imp_cscore | RiskScore | axio_score |
|-------|---------------|--------------|--------------|----------------|--------------|--------------|--------------|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 |
| mean | 496819.831400 | 65853.355259 | 20.962621 | 0.329419 | 662.548800 | 670.042869 | 0.393211 |
| std | 287391.314157 | 31093.369592 | 18.841121 | 0.223299 | 90.549985 | 89.965854 | 0.288243 |
| min | 244.000000 | 2.054543 | -2.140206 | 0.001781 | 500.000000 | 324.436647 | -0.000052 |
| 25% | 245172.500000 | 39165.786086 | 6.150577 | 0.156965 | 600.000000 | 609.231181 | 0.139424 |
| 50% | 495734.000000 | 76903.628763 | 11.913366 | 0.272263 | 655.000000 | 669.493442 | 0.337841 |

| | | | | | | | |
|-----|---------------|---------------|-----------|----------|-------------|--------------|----------|
| 75% | 745475.250000 | 91082.514000 | 32.236904 | pre_prob | 727mp_score | 730RiskScore | 0.6248 |
| max | 999870.000000 | 150538.809704 | 81.759632 | 1.144357 | 849.000000 | 1004.497869 | 1.000000 |

In [160]:

```
cat_df = df_updated.select_dtypes(include=['object']).copy()
```

In [162]:

```
cat_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 3 columns):
demographic_slice    10000 non-null object
country_reg          10000 non-null object
ad_exp               10000 non-null object
dtypes: object(3)
memory usage: 234.5+ KB
```

In [163]:

```
cat_df.head()
```

Out[163]:

| | demographic_slice | country_reg | ad_exp |
|---|-------------------|-------------|--------|
| 0 | AX03efs | W | N |
| 1 | AX03efs | E | N |
| 2 | AX03efs | W | Y |
| 3 | AX03efs | E | Y |
| 4 | AX03efs | W | N |

In [164]:

```
df_updated['demographic_slice']= df_updated['demographic_slice'].astype('category')
```

In [165]:

```
#Label Encoding
df_updated['demographic_slice']= df_updated['demographic_slice'].cat.codes
```

In []:

```
#similary for rest of the other
```

In [166]:

```
df_updated['country_reg']= df_updated['country_reg'].astype('category')
df_updated['ad_exp']= df_updated['ad_exp'].astype('category')
df_updated['card_offer']= df_updated['card_offer'].astype('category')
```

In [167]:

```
#Label Encoding
df_updated['country_reg']= df_updated['country_reg'].cat.codes
df_updated['ad_exp']= df_updated['ad_exp'].cat.codes
df_updated['card_offer']= df_updated['card_offer'].cat.codes
```

In [168]:

```
df_updated.head(10)
```

Out[168]:

| | customer_id | demographic_slice | country_reg | ad_exp | est_income | hold_bal | pref_cust_prob | imp_cscore | RiskScore |
|---|-------------|-------------------|-------------|--------|--------------|-----------|----------------|------------|-----------|
| 0 | 713782 | 0 | 1 | 0 | 33407.901749 | 3.000000 | 0.531112 | 619 | 503.24902 |
| 1 | 515901 | 0 | 0 | 0 | 19927.533533 | 20.257927 | 0.297439 | 527 | 820.10814 |
| 2 | 95166 | 0 | 1 | 1 | 51222.470997 | 4.000000 | 0.018463 | 606 | 586.60579 |
| 3 | 425557 | 0 | 0 | 1 | 67211.587467 | 18.653631 | 0.089344 | 585 | 634.70198 |
| 4 | 624581 | 0 | 1 | 0 | 20093.342158 | 4.000000 | 0.094948 | 567 | 631.94997 |
| 5 | 721691 | 0 | 0 | 0 | 73896.096129 | 12.906641 | 0.656848 | 560 | 809.33396 |
| 6 | 269858 | 0 | 1 | 0 | 73609.404135 | 3.000000 | 0.137818 | 620 | 697.30816 |
| 7 | 219196 | 0 | 1 | 0 | 57619.668582 | 9.000000 | 0.367879 | 658 | 668.07547 |
| 8 | 413020 | 0 | 1 | 1 | 49282.620299 | 0.000000 | 0.182079 | 519 | 656.11159 |
| 9 | 174424 | 0 | 1 | 0 | 57173.061392 | 6.000000 | 0.288242 | 645 | 547.76711 |

In [169]:

```
df_updated.columns
```

Out[169]:

```
Index(['customer_id', 'demographic_slice', 'country_reg', 'ad_exp',  
      'est_income', 'hold_bal', 'pref_cust_prob', 'imp_cscore', 'RiskScore',  
      'axio_score', 'card_offer'],  
      dtype='object')
```

In [170]:

```
offer= df_updated['card_offer']  
users= df_updated['customer_id']  
  
df_final= df_updated.drop(columns=['card_offer','customer_id'])
```

In [172]:

```
#splitting the data in training and test DataSets  
  
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(df_final, offer, test_size=0.2, random_state=0  
)
```

In [173]:

```
#feature scaling  
  
from sklearn.preprocessing import StandardScaler  
sc_X = StandardScaler()  
X_train2 = pd.DataFrame(sc_X.fit_transform(X_train))  
X_test2 = pd.DataFrame(sc_X.fit_transform(X_test))  
  
X_train2.columns= X_train.columns.values  
X_test2.columns= X_test.columns.values  
X_train2.index= X_train.index.values  
X_test2.index= X_test.index.values  
  
X_train= X_train2  
X_test= X_test2  
  
#X_test= np.squeeze(sc_X.transform(X_test.values.reshape(-1, 1)))
```

In []:

```
#ready for modelling
```

```
#Logistic Reg
```

In [174]:

```
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state=0)
classifier.fit(X_train, y_train)
```

Out[174]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=0, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

In [175]:

```
#predicting Test Set
```

```
y_pred_LR= classifier.predict(X_test)
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score , recall_score
acc= accuracy_score (y_test, y_pred_LR)
prec= precision_score(y_test,y_pred_LR)
rec= recall_score (y_test, y_pred_LR)
f1= f1_score(y_test,y_pred_LR)
```

In [176]:

```
pd.DataFrame([['Log Regr', acc, prec, rec, f1]], columns= ['Model','Accuracy','Precision','Recall',
'F1'])
```

Out[176]:

| | Model | Accuracy | Precision | Recall | F1 |
|---|----------|----------|-----------|----------|----------|
| 0 | Log Regr | 0.9665 | 0.905724 | 0.873377 | 0.889256 |

In []:

```
#lets check with different model
```

In [177]:

```
from sklearn.svm import SVC
classifier = SVC(random_state=0, kernel='linear')
classifier.fit(X_train, y_train)

#predicting Test Set

y_pred_SVC= classifier.predict(X_test)
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score , recall_score
acc= accuracy_score (y_test, y_pred_SVC)
prec= precision_score(y_test,y_pred_SVC)
rec= recall_score (y_test, y_pred_SVC)
f1= f1_score(y_test,y_pred_SVC)

pd.DataFrame([['SVM', acc, prec, rec, f1]], columns= ['Model','Accuracy','Precision','Recall','F1'])
```

Out[177]:

| | Model | Accuracy | Precision | Recall | F1 |
|---|-------|----------|-----------|----------|----------|
| 0 | SVM | 0.966 | 0.902685 | 0.873377 | 0.887789 |

Conculsuiion from above results:

1. RBF kernel worked better than linear kernel but not good than that of Log Regr

In [178]:

```
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(random_state=0, n_estimators=100, criterion = 'entropy')
classifier.fit(X_train, y_train)

#predicting Test Set

y_pred_RC= classifier.predict(X_test)
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score , recall_score
acc= accuracy_score (y_test, y_pred_RC)
prec= precision_score(y_test,y_pred_RC)
rec= recall_score (y_test, y_pred_RC)
f1= f1_score(y_test,y_pred_RC)

pd.DataFrame([['Random Forest with 100 trees', acc, prec, rec, f1]], columns= ['Model','Accuracy',
'Precision','Recall','F1'])
#print(svm_results)
```

Out[178]:

| | Model | Accuracy | Precision | Recall | F1 |
|---|------------------------------|----------|-----------|----------|----------|
| 0 | Random Forest with 100 trees | 0.9755 | 0.94198 | 0.896104 | 0.918469 |

Final Conclusion

RFC performed best among the rest of the classifiers with an accuracy of 97.55 percent and F1 score 0.918 on the training

first dataset named "Test1.csv" , please note , no dataset ds1 & ds4 was provided in mail communication"

I will proceed with the RFC model for the test2.csv data